


The Smol Training Playbook: The Secrets to Building World-Class LLMs



A practical journey through the challenges, decisions, and messy reality behind training state-of-the-art language models

AUTHORS	AFFILIATION	PUBLISHED
Loubna Ben Allal , Lewis Tunstall , Nouamane Tazi , Elie Bakouch , Ed Beeching , Carlos Miguel Patiño , Clémentine Fourrier , Thibaud Frere , Anton Lozhkov , Colin Raffel , Leandro von Werra , Thomas Wolf	Hugging Face	Oct. 30, 2025

Introduction

What does it actually take to train a high-performance LLM today?

Published research makes it look straightforward: strategic architecture choices, carefully curated datasets, and sufficient compute. The results are polished, the ablations are structured and clean. Every decision seems obvious in hindsight. But those reports only show what worked and apply a bit of rosy retrospection—they don't capture the 2 a.m. dataloader debugging sessions, the loss spikes, or the subtle tensor parallelism bug (see later!) that quietly sabotages your training. The reality is messier, more iterative, and full of